

ECHOSPHERE 2026 · ROUND IV · PROJECT SUBMISSION & EVALUATION

Close

A voice AI sales agent that listens, adapts, and actually closes.

Deployed. Tested. 38 automated tests passing. 3 real integrations.

Team CtrlPlusProfit

TEAM DETAILS

Who's building this

Name	Role	Email
Swati Aggrawal	Technical Lead — AI & Voice Engineering	swatiagg357@gmail.com
Aashi Aggarwal	Sales Strategy & Conversation Design Lead	aashiaggarwal346@gmail.com

Swati built and shipped the full backend — voice pipeline integration, all tools, real Calendar/Sheets/email integrations, and 38 automated tests.

Aashi shaped the conversation design — objection handling, qualification flow, and the client-facing tone across every automated email and confirmation.

Note: emails are placeholders — replace with real contact addresses before submission.

PROBLEM STATEMENT

Why scripts don't sell

Build a real-time voice AI sales agent that conducts a complete customer qualification and sales conversation without depending on a fixed script — one that handles natural turn-taking and interruptions, adapts to objections and changing requirements, retrieves live product and pricing information, and drives every call toward a meaningful outcome such as a booked meeting or a qualified lead.

Broken trust

Scripted replies feel robotic the moment a customer goes off-book

Missed signals

Real buying intent and objections get steamrolled by fixed flows

Lost revenue

Leads that needed one adaptive answer get qualified out instead of in

PROPOSED SOLUTION

Meet Close — now live, not just conceptual

Close is a real-time voice sales agent built on Agora's Conversational AI Engine. Agora's own managed LLM handles reasoning, memory, and turn-taking natively — Close's job is the tool layer beneath it: real pricing logic, real retrieval, and real business actions (CRM, email) triggered directly from the conversation.

Adaptive, not scripted

Goal-oriented reasoning
replaces fixed dialogue trees

Always retrieving

Live pricing and service info
pulled mid-call, never
guessed


Never forgets

Session memory holds every
detail across the whole call

Takes real action

Calendar events, CRM rows,
and emails — not just log
lines

System architecture



WHAT'S ACTUALLY REAL

No mockups — real external side effects

- **Live Dashboard**

Real events created during the conversation

- **Google Sheets**

Every lead appended as a live row — a genuine, viewable CRM, not a local file

- **Resend Email**

Internal team notifications AND client-facing confirmations, both real sends

Retrieval is keyword-scored over a small document set for this round, not full embeddings-based RAG — it correctly grounds every answer in a real document, using a simpler technique than the label "RAG" usually implies.

CONVERSATION INTELLIGENCE

It talks like a person, not a menu

Natural turn-taking

Full-duplex audio lets the customer interrupt and be heard immediately

Session memory

A requirement changed mid-call is tracked correctly, never reverts to the old value

Grounded answers

Competitor and guarantee questions are answered from real documents, not guesses

REAL EXCHANGE, LIVE-TESTED

Customer: *"I need 20 finance professionals, what would that cost?"*

Close: *"\$21,600 — that's \$1,080 per candidate after a volume discount. Want to talk timeline?"*

Customer: *"Actually, let's change that to 50."*

Close: *"Got it, updating that — for 50, you'd be at \$960 per candidate, \$48,000 total."*

Customer: *"Just to confirm, how many did I say I needed?"*

Close: *"Fifty finance professionals."*

OBJECTION HANDLING & QUALIFICATION

Turning pushback into progress

Objection Type	Detection Signal	Close's Strategy
Pricing pushback	Cost mentioned relative to competitor or budget	Reframes around value delivered per seat, offers tiered options
Competitor comparison	Named competitor product or feature	Pulls live differentiators from the retrieval layer, stays factual
Trust / credibility	Hesitation, vague stalling language	Offers proof points or a human hand-off
Timing / budget	"Not right now" or "need approval"	Locks in a qualified follow-up instead of losing the lead

QUALIFICATION FRAMEWORK — BANT, ASKED CONVERSATIONALLY



Two layers, matched to what each thing controls

Automated

38 tests

Pricing math, discount tiers, retrieval scoring, email validation, graceful degradation across every integration. Runs in under a second, no live call needed.

Manual, scripted

8 scenarios

Memory, interruption, objection handling, retrieval grounding, escalation, and the consolidated booking flow — run against the live agent through Agora.

Live dashboard

Real-time

An ops view reading off the same logs every tool writes to — quotes given, leads captured, meetings booked, escalations, updating as calls happen.

Why two layers, not one: conversational reasoning lives inside Agora's managed LLM, not our code — it can't be unit tested. We test what we control (tool logic) automatically, and what Agora controls (conversation behavior) through live scripted calls.

KNOWN LIMITATIONS

Stated plainly, not discovered live

Voice ASR on spelled-out contact details

Emails and phone numbers spoken letter-by-letter are a hard case for any speech recognition. Mitigated with a mandatory read-back-and-confirm step before the agent proceeds, plus server-side normalization — not fully solvable, but caught before it reaches a human.

Retrieval technique

Keyword-scored matching over a small document set, not embeddings-based RAG. Functionally grounds every answer correctly, but it's a simpler technique than the label usually implies — worth being precise about if asked.

The booking consolidation itself

Root-caused a real platform-level reliability issue rather than shipping around it silently — the architecture on slide 5 is the direct result of that diagnosis, not a workaround we're hiding.

What makes this team's take different

- The system is scoped end-to-end and actually deployed — voice pipeline, tool layer, and three real external integrations all work together right now, not just on paper.
- Pairing a technical builder with a client-facing sales lead means objection handling and qualification are grounded in how real conversations sound, not just how they look on a whiteboard.
- Every architecture decision in this deck — dropping the custom LLM adapter, consolidating booking into create_lead — came from diagnosing a real problem in live testing, not from a plan made in advance.